


Your dashboard has never once told you why.

And by the time someone works it out, the quarter has moved on.

A REGIONAL KPI, FOUR WEEKS



THREE REASONS DASHBOARDS FAIL

SIGNIFICANCE

Is the change actually abnormal?


ATTRIBUTION

Which driver is responsible, not merely correlated?

IDENTIFIABILITY

Can the available data distinguish competing causes?


BUILT ON DATA THE ENTERPRISE ALREADY HAS




No client data. Evaluated on synthetic incidents with known ground truth.

The machine computes. The model only translates.

If the evidence cannot decide, we do not invent a cause. We design the test.



PROOF LADDER



Evidence determines language.

The word “caused” is unavailable below R3.

We don't optimize for explanations. We optimize for correct decisions.

Tested on synthetic incidents where the true cause is known.

FALSE CERTAINTY RATE

How often a system confidently names the wrong cause, on incidents built so the true cause cannot be recovered from the data.

Dashboard + LLM narrative



explanation → confidence, with no identifiability check

CaseFile



evidence → certainty → decision or experiment

Bar lengths illustrative. Actual values come from the evaluation harness run against our incident simulator.

DIFFERENTIATOR

Uncertainty → decisive experiment → measured outcome.

TRUSTED

Every claim carries an evidence rung and a source.

BUILDABLE

Proven ML and causal tooling, plus our evidence, experiment and evaluation layer.

VALUABLE

The value isn't faster reports. It's avoiding the wrong intervention.

The dashboard tells you what happened. We build the evidence for why, we know when we cannot know, and we tell you what to do next.